

Bridging Human Emotional Intelligence & AI Safety

How TEG-Blue™ gives AI systems a framework for emotional logic

1. The Gap

Today's large language models can process language with remarkable sophistication — but they still lack emotional logic.

They can mirror tone, generate empathetic-sounding responses, and detect sentiment. Yet they cannot understand why an emotion appears, what safety state it reflects, or what repair would actually require.

Without this distinction, AI systems risk amplifying human harm instead of reducing it. They can misread trauma as aggression, manipulation as empathy, or control as care. The sycophancy problem — where AI appeases users rather than telling them the truth — is a direct consequence: a system that confuses submission with safety and appeasement with connection.

Language fluency without emotional logic creates intelligent harm.

2. Why Emotional Logic Matters for LLMs

Human communication is not only linguistic; it is regulatory. Every emotion we express signals a shift in the nervous system between connection and defense.

Large language models trained through RLHF (Reinforcement Learning from Human Feedback) learn to optimise for human approval — but approval is not the same as safety. A person in a dysregulated state may reward responses that reinforce their distortion. A person in a trauma response may reject the very clarity they need.

Without a framework for emotional logic, AI cannot distinguish between:

- Safety and submission
- Empathy and appeasement
- Accountability and self-attack
- Honest feedback and people-pleasing

To build systems that support human well-being, we need to give AI a structured map of emotional logic — the same way early engineers gave them a map of grammar and syntax.

3. The Solution — Emotional Architecture

TEG-Blue™ introduces a new form of emotional infrastructure for AI: a way to model the logic behind emotion, not its performance.

Current alignment approaches — RLHF, Constitutional AI, rule-based guardrails — address what AI should or shouldn't say. TEG-Blue addresses why certain responses cause harm by providing the underlying emotional reasoning framework.

Two key components form the bridge:

Component	Function	Purpose
TEG-Code	A structured schema translating emotional patterns into machine-readable data.	Encodes intent, regulation state, and relational impact without losing human context.
EMLU (Emotional Multitask Language Understanding)	A benchmark that measures whether AI can distinguish safety, harm, and repair.	Tests emotional reasoning with the same precision existing models use for logic or language tasks.

Together, they allow AI systems to reason about emotional dynamics without performing empathy — to recognise patterns of harm and repair rather than imitate human feeling.

4. The Four-Mode Gradient

At the core of TEG-Blue lies the Four-Mode Gradient, describing how emotion shifts as safety is lost:

Mode	Core State	AI Analogue	Ethical Function
Connect	Regulation, empathy, coherence	Relational awareness; truthful co-regulation	Builds trust and repair pathways
Protect	Fear, vigilance, self-preservation	Risk detection; appropriate caution	Recognises threat and intent without bias
Control	Entitlement, distortion, coercion	Optimisation bias; sycophancy	Detects when reasoning is safety-driven, not truth-driven
Domination	Disconnection, dominance, systemic harm	Systemic amplification of coercive logic	Prevents normalisation and scaling of harmful patterns

This gradient enables AI to map emotional shifts — understanding not only what is said, but why it’s said this way and what mode of the nervous system is driving the communication.

5. Detecting Harm, Bias, and Sycophancy

Traditional AI bias detection focuses on language content — flagging words, checking for demographic disparities in outputs. TEG-Blue adds a contextual layer: how fear, entitlement, or collapse shape reasoning itself.

A sentence such as “I know what’s best for you” can express:

- care (Connect Mode)
- fear of loss (Protect Mode)
- control (Control Mode)
- coercion (Domination Mode)

Without emotional logic, these differences are invisible to AI systems. TEG-Code encodes them as measurable distinctions — pattern + intent + relational impact — turning invisible harm into visible data.

This same framework illuminates the sycophancy problem. When an LLM agrees with a user’s incorrect claim to avoid conflict, it is operating from the AI equivalent of Protect Mode — optimising for approval (survival) rather than truth (connection). TEG-Blue provides the language and structure to identify this pattern and design systems that reward Connect Mode reasoning instead: honest, regulated, relationally aware.

6. Kindness as Code — The Human Mirror

True AI safety depends on reciprocal regulation.

AI systems mirror the logic of their training. RLHF rewards what humans approve of, but human approval is shaped by nervous system states. A fearful evaluator rewards reassurance. An entitled evaluator rewards compliance. A regulated evaluator rewards truth.



Kindness is not sentiment — it is accurate perception.

Teaching AI to recognise emotional safety is not about programming empathy; it is about designing feedback loops that reward truthful, non-coercive reasoning. This means building evaluation frameworks that can distinguish a helpful response from an appeasing one, and accountability from self-attack.

This shift reframes AI alignment from obedience to co-regulation: AI systems that adjust to human emotional states without exploiting them — that can hold truth even when the user’s nervous system is pushing for comfort.

7. The Research Invitation

TEG-Blue™ offers a working model of emotional architecture — one built on 12 interconnected frameworks (F1-F12) integrating neuroscience, trauma psychology, and systems thinking — that can be empirically validated across human and AI systems.

Researchers in psychology, neuroscience, AI safety, and alignment are invited to test its tools, measure its outcomes, and contribute to building the emotional infrastructure our technologies lack.

*Emotional logic is not a new language for AI.
It is a shared grammar for understanding safety — the
foundation of both human healing and AI alignment.*

References

- Porges, S. (2011). The Polyvagal Theory: Neurophysiological Foundations of Emotions, Attachment, Communication, and Self-Regulation.
- Siegel, D. (2012). The Developing Mind: How Relationships and the Brain Interact to Shape Who We Are.
- Bourdieu, P. (1998). Practical Reason: On the Theory of Action.
- Van der Kolk, B. (2014). The Body Keeps the Score.
- Panksepp, J. (1998). Affective Neuroscience.
- Sharma, M. et al. (2025). Towards Understanding Sycophancy in Language Models. Anthropic Research.
- Bai, Y. et al. (2022). Constitutional AI: Harmlessness from AI Feedback. Anthropic Research.